

Lab 0

Nicholas Storti

June 13, 2026

1 GPU Output

There is 1 device supporting CUDA

```
Device 0: "Tesla V100-PCIE-16GB"
Major revision number:          7
Minor revision number:         0
Total amount of global memory: 4043440128 bytes
Number of multiprocessors:     80
Number of cores:               640
Total amount of constant memory: 65536 bytes
Total amount of shared memory per block: 49152 bytes
Total number of registers available per block: 65536
Warp size:                     32
Maximum number of threads per block: 1024
Maximum sizes of each dimension of a block: 1024 x 1024 x 64
Maximum sizes of each dimension of a grid: 2147483647 x 65535 x 65535
Maximum memory pitch:          2147483647 bytes
Texture alignment:             512 bytes
Clock rate:                    1.38 GHz
Concurrent copy and execution: Yes
```

TEST PASSED

2 Instructional Output

There is 1 device supporting CUDA

```
Device 0: "NVIDIA L40S"
Major revision number:          8
Minor revision number:         9
Total amount of global memory: 432537600 bytes
Number of multiprocessors:     142
Number of cores:               1136
```

Total amount of constant memory:	65536 bytes
Total amount of shared memory per block:	49152 bytes
Total number of registers available per block:	65536
Warp size:	32
Maximum number of threads per block:	1024
Maximum sizes of each dimension of a block:	1024 x 1024 x 64
Maximum sizes of each dimension of a grid:	2147483647 x 65535 x 65535
Maximum memory pitch:	2147483647 bytes
Texture alignment:	512 bytes
Clock rate:	2.52 GHz
Concurrent copy and execution:	Yes

TEST PASSED